

Change Summary

Every change made to the application since the original codebase, with the reason for it. Type is **Fix** (the original was defective), **Changed** (the original worked but had to change for a stated constraint), or **New** (a capability that did not exist).

Status as of 2026-09-14. Detail and test names for each item are in PROGRESS.md; plain-language workstream status is in WORKSTREAMS.md.

1. Speaker identification (reported challenge 1)

Type	Area	Original	Now	Rationale
Fix	Speaker correction pass	An unmeasured speaker label scored 1.0, so the correction test could never pass for "Speaker N" labels	Unmeasured label scores 0.0; accept 0.82 and minimum 0.20 thresholds are settings, defaults unchanged	The labels most likely to be wrong were the only ones the correction could not touch
Fix	Voice enrolment errors	Every failure said "Could not extract embeddings. Please re-record."	Each sample classed as too short, too quiet or empty (user can fix) or server fault (re-recording will not help)	The live error was a missing server library; asking the user to re-record could never work
Fix	Audio decoding	torchcodec, which needs separate FFmpeg libraries, failed on Windows	librosa and soundfile decode audio	Enrolment and overlap detection failed outright

2. Short and medium ROM losing information (reported challenge 2)

Type	Area	Original	Now	Rationale
Fix	Metadata after condensing	Owner, dates, figures and terms reattached by list position after the model reordered points; dates, numbers and terms emptied	Model returns source point ids; metadata merged from every source point	This is the information loss ADA reported
New	Coverage record	Dropped points vanished silently	Each version records which points were kept, merged and dropped	A reviewer can see and restore what was cut
New	Importance scoring	No notion of which points matter	Seven-signal score shown to the model when condensing	Condensing keeps what matters instead of guessing
Fix	Empty AI instructions	Four prompt keys sent an empty prompt to the model	All wired to real templates; an unknown or empty key raises an error	Model output for those steps was unguided
Fix	Stray import	Condensing service imported an internal transformers test module	Removed	Could crash on a different library version

3. Action points (reported challenge 3)

Type	Area	Original	Now	Rationale
Fix	Action owners	Unassigned actions given to every speaker on the point	Owner only from evidence, with its source labelled: named by model, extracted, named in task, inferred from sole speaker, or unassigned	Guessed owners were presented as fact
Fix	PDF export	Speaker summary section used values before they were built	Values built first	Export crashed
Fix	Training variants	New prompt variants could reuse an existing id and number	Unique id and next free number	One variant silently overwrote another

4. Offline operation (no internet at runtime)

Type	Area	Original	Now	Rationale
Fix	Diarization telemetry	pyannote sent usage traces to its vendor server by default; 78 attempts on one 15 s recording	Forced off, with Hugging Face, ChromaDB and OpenTelemetry opt-outs, before any model loads	Outbound data from a defence installation
Fix	Search database telemetry	ChromaDB sent usage events	Disabled	Same
Fix	Fonts	Loaded from Google Fonts	34 font files bundled with the app	Blank or slow first paint offline, plus outbound requests
Fix	Build dependency	lovable-tagger, an external service plugin, in the frontend build	Removed	External dependency, and it blocked installation
Changed	API docs pages	/docs and /redoc on by default, loading scripts from public CDNs	Off unless a developer enables them	Outbound requests and a broken page offline
Changed	Desktop app loading	Electron loaded files from disk with web security relaxed	Private app:// protocol, path traversal blocked, web security on	Relative assets broke and security was weakened
New	Offline test	None	Real recording processed with all outbound connections blocked; zero attempts	Proves the offline claim rather than assuming it

5. Model choice (no Chinese-origin models)

Type	Area	Original	Now	Rationale
Changed	Language model order	Qwen and DeepSeek preferred	Phi-4 (Microsoft), then Mistral Small, Llama 3.1, others; Qwen and DeepSeek removed	Stated sourcing constraint; Phi-4 rated best of the permitted options
Changed	Text embedding model	Qwen embedding model	mxbai-embed-large-v1 (Mixedbread, Germany); picker lists only permitted models	Same constraint
Fix	Embedding pooling	Pooling method assumed	Read from each model's own configuration	Wrong pooling degrades document search
Changed	Interface labels	"Powered by Qwen AI"	"Local AI, runs entirely on this machine"	Label was inaccurate
New	Settings migration	Existing accounts kept Qwen settings and the AI switched off	Old values converted once at startup	Your account's settings were wrong after the change
Open	OCR component	RapidOCR uses Baidu models	Not yet replaced	Remaining sourcing issue

6. Security and access control

Type	Area	Original	Now	Rationale
New	Administrator role	Dashboard, logs and maintenance open to any logged-in user	Role per user, first account is admin, admin-only routes enforced including live updates	Anyone could read logs and run maintenance
Fix	Login token key	Every install shipped the same known secret key	Unique key generated per install and kept out of Git	Anyone with the code could forge logins
Fix	Overlap detection endpoint	No login required	Login required	Unauthenticated access to audio processing
Changed	Password hashing	passlib, unmaintained and broken with current bcrypt	bcrypt directly; passwords over 72 bytes rejected instead of silently truncated	Registration and login failed
Fix	Leaked log file	Diagnostics export containing logs committed to GitHub	Removed from repository and ignored	Log contents in version control

7. Reliability and performance

Type	Area	Original	Now	Rationale
Fix	Stalled language model	Requests had no timeout; a stalled model froze processing forever	Streaming with a stall timeout; truncated answers rejected	Found when the evaluation hung indefinitely
Fix	"Database is locked"	Default SQLite settings	Write-ahead logging, 10 s wait on locks	Long processing blocked the interface
Fix	Launch folder	Training data and diagnostics export used the current folder	Fixed application paths; export uses a temporary file deleted after sending	Launching from elsewhere read and wrote the wrong place
New	Run metrics	No timing or usage data	Per-run stage timing, model loads and exact token counts stored	Needed to prove any speed-up
Fix	Runtime crash defects	Settings page referenced a shadowed variable; several missing imports	Fixed; crash-only lint added	Found by lint, would fail at runtime

8. Installation, build and testing

Type	Area	Original	Now	Rationale
Fix	Python requirements	Could not install: conflicting torch and bcrypt versions	Torch 2.8 pinned, GPU build chosen automatically	A fresh clone could not be installed
New	One-command setup	Manual steps	setup.py creates the environment, installs, fetches models, pulls Phi-4, and checks the result	A clone should work upfront
New	Automatic checks on GitHub	None	Backend tests, crash lint and frontend build on every push	Catches breakage before a demo
New	Golden dataset harness	No way to measure quality	Scores word error, speaker error, information kept, must-keep points, action precision and recall, owner accuracy; 3 synthetic cases; draft-reference helper	Every quality claim needs a number; real recordings still needed
Fix	Test isolation	Tests wrote into your real database, changed tracked files and called your live model	Temporary database and files, network blocked in tests	Test runs polluted real data; suite dropped from 26 to about 4 minutes
New	Tests	Original suite	726 passing, about 25 new test files	Each fix above has a test that fails if it regresses
Changed	Repository hygiene	Docs ignored; no rules for secrets or audio	Docs tracked; keys, secrets, audio, local datasets and virtual environment ignored	Keep secrets and meeting audio out of Git

9. Documentation

Type	Document	Purpose
New	observations/2026-09-12-parameter-audit.md	Every tunable parameter and findings against the three challenges
New	REVAMP-PLAN.md	End-to-end plan with rationale for each change
New	IMPLEMENTATION-AND-VALIDATION.md	Execution order and how each goal is proven
New	OBSERVABILITY-AND-AUDIT.md	Audit trail, traceability and weekly leadership report design
New	WORKSTREAMS.md	Plain-language status of all 42 workstreams
New	PROGRESS.md	Detailed record of what was built and tested
New	CHANGE-SUMMARY.md	This table
New	SETUP.md, backend/eval/README.md	Installation and golden dataset instructions